

# 4\_25

Tony

2026-03-31

## Exercise 4.25: Housing Price Analysis in Collegetown

```
# Load necessary Libraries
library(POE5Rdata)
library(tseries)
library(knitr)
library(ggplot2)

# Load data
data(collegetown)
summary(collegetown)
```

```
##      price      sqft      age      pool
## Min.   : 50.0   Min.   :10.00  Min.   : 1.00  Min.   :0.000
## 1st Qu.:157.9   1st Qu.:20.56  1st Qu.: 7.00  1st Qu.:0.000
## Median :199.9   Median :25.55  Median : 9.00  Median :0.000
## Mean   :250.2   Mean   :27.28  Mean   : 8.19  Mean   :0.064
## 3rd Qu.:279.0   3rd Qu.:31.20  3rd Qu.:10.00  3rd Qu.:0.000
## Max.   :1370.0   Max.   :91.67  Max.   :11.00  Max.   :1.000
## fireplace      close      twostory      occupied
## Min.   :0.000   Min.   :0.000   Min.   :0.000   Min.   :0.00
## 1st Qu.:0.000   1st Qu.:0.000   1st Qu.:0.000   1st Qu.:0.00
## Median :1.000   Median :0.000   Median :0.000   Median :1.00
## Mean   :0.546   Mean   :0.378   Mean   :0.138   Mean   :0.56
## 3rd Qu.:1.000   3rd Qu.:1.000   3rd Qu.:0.000   3rd Qu.:1.00
## Max.   :1.000   Max.   :1.000   Max.   :1.000   Max.   :1.00
```

```
# Calculate sample means for elasticity and slope calculations
avg_sqft <- mean(collegetown$sqft)
avg_price <- mean(collegetown$price)
```

### (a) Log-Linear Model

```
mod_a <- lm(log(price) ~ sqft, data = collegetown)
summary(mod_a)
```

```
##
## Call:
## lm(formula = log(price) ~ sqft, data = collegetown)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -1.32528 -0.19199  0.00122  0.21678  0.86609
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  4.393866   0.043288  101.50  <2e-16 ***
## sqft         0.036044   0.001486   24.26  <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.34 on 498 degrees of freedom
## Multiple R-squared:  0.5417, Adjusted R-squared:  0.5408
## F-statistic: 588.7 on 1 and 498 DF,  p-value: < 2.2e-16
```

```
b2 <- coef(mod_a)["sqft"]

# Calculating slope and elasticity at sample means
# Slope:  $dP/dx = \beta_2 * P$ 
# Elasticity:  $\epsilon = \beta_2 * x$ 
slope_a <- b2 * avg_price
elasticity_a <- b2 * avg_sqft
```

**Interpretation:** The estimated coefficient  $\hat{\beta}_2$  is approximately 0.036. This suggests that for every additional 100 square feet, the house price is expected to increase by approximately 3.6%. At the sample means, the marginal effect (slope) is 9.02 and the elasticity is 0.98.

## (b) Log-Log Model

```
mod_b <- lm(log(price) ~ log(sqft), data = collegetown)
summary(mod_b)
```

```
##
## Call:
## lm(formula = log(price) ~ log(sqft), data = collegetown)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -1.36864 -0.20849  0.00487  0.23204  1.21099
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  2.04965    0.15797   12.97  <2e-16 ***
## log(sqft)    1.02483    0.04839   21.18  <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.3643 on 498 degrees of freedom
## Multiple R-squared:  0.4738, Adjusted R-squared:  0.4728
## F-statistic: 448.5 on 1 and 498 DF,  p-value: < 2.2e-16
```

```
a2 <- coef(mod_b)["log(sqft)"]

# Slope and Elasticity
# Slope:  $dP/dx = \alpha_2 * (P/x)$ 
# Elasticity:  $\alpha_2$ 
slope_b <- a2 * (avg_price / avg_sqft)
elasticity_b <- a2
```

**Interpretation:** The coefficient  $\hat{\alpha}_2$  is 1.02, which represents a constant elasticity. A 1% increase in square footage is associated with a 1.02% increase in house price. At the sample means, the slope is 9.4.

## (c) Comparison of Goodness-of-Fit

```
mod_c <- lm(price ~ sqft, data = collegetown)
summary(mod_c)
```

```
##
## Call:
## lm(formula = price ~ sqft, data = collegetown)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -316.93  -58.90   -3.81   47.94  477.05
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) -115.4236    13.0882  -8.819  <2e-16 ***
## sqft         13.4029     0.4492  29.840  <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 102.8 on 498 degrees of freedom
## Multiple R-squared:  0.6413, Adjusted R-squared:  0.6406
## F-statistic: 890.4 on 1 and 498 DF,  p-value: < 2.2e-16
```


```
# Generalized R-squared for Log Models
y_true <- collegetown$price
gen_r2_a <- cor(y_true, exp(predict(mod_a)))^2
gen_r2_b <- cor(y_true, exp(predict(mod_b)))^2
r2_c <- summary(mod_c)$r.squared
```

**Answer:** Comparing the generalized  $R^2$ , the Log-Linear model 0.6622 and Log-Log model 0.6445 both provide a better fit to the data than the simple linear model 0.6413.

## (d) Normality of Residuals


```
hist(residuals(mod_a), breaks = 20, col = "skyblue", main = "Log-Linear Res")
```

### Log-Linear Res




```
hist(residuals(mod_b), breaks = 20, col = "lightgreen", main = "Log-Log Res")
```

### Log-Log Res



```
hist(residuals(mod_c), breaks = 20, col = "salmon", main = "Linear Res")
```




```
# Jarque-Bera Tests
jb_a <- jarque.bera.test(residuals(mod_a))
jb_b <- jarque.bera.test(residuals(mod_b))
jb_c <- jarque.bera.test(residuals(mod_c))
```

**Answer:** The Jarque-Bera test p-values are extremely low for all models, especially for the linear model ( $p \approx 0$ ). This indicates that the residuals are not normally distributed. However, the log-transformed models show less extreme skewness in the histograms.

## (e) Residual Patterns and Heteroscedasticity


```
plot(collegetown$sqft, residuals(mod_a), main = "Log-Linear: Res vs SQFT")
abline(h = 0, col = "red")
```

### Log-Linear: Res vs SQFT




```
plot(collegetown$sqft, residuals(mod_b), main = "Log-Log: Res vs SQFT")  
abline(h = 0, col = "red")
```

### Log-Log: Res vs SQFT



```
plot(collegetown$sqft, residuals(mod_c), main = "Linear: Res vs SQFT")
abline(h = 0, col = "red")
```



**Answer:** The linear model's residuals show a distinct fan shape, indicating heteroscedasticity. The log-transformed models (a) and (b) help stabilize this variance, though some patterns may still remain.



## (f)(g) Prediction for a 2700 sqft House

```
new_data <- data.frame(sqft = 27) # 2700 sqft corresponds to sqft = 27 in this dataset

# Predictions (Exp-transformed for Log models)
pred_a_log <- predict(mod_a, new_data, interval = "prediction", level = 0.95)
pred_a <- exp(pred_a_log)

pred_b_log <- predict(mod_b, new_data, interval = "prediction", level = 0.95)
pred_b <- exp(pred_b_log)

pred_c <- predict(mod_c, new_data, interval = "prediction", level = 0.95)

# Summary Table
results_comparison <- data.frame(
  Model = c("Log-Linear", "Log-Log", "Linear"),
  R_squared = c(gen_r2_a, gen_r2_b, r2_c),
  JB_p_value = c(jb_a$p.value, jb_b$p.value, jb_c$p.value),
  Pred_Price_2700 = c(pred_a[1], pred_b[1], pred_c[1])
)

kable(results_comparison, caption = "Comparison of Model Results and Predictions", digits = 4)
```

Comparison of Model Results and Predictions

| Model      | R_squared | JB_p_value | Pred_Price_2700 |
|------------|-----------|------------|-----------------|
| Log-Linear | 0.6622    | 0e+00      | 214.2336        |
| Log-Log    | 0.6445    | 8e-04      | 227.5386        |
| Linear     | 0.6413    | 0e+00      | 246.4557        |

## (h) Conclusion on Functional Form

**Answer:** Based on the analysis, the Log-Log model (b) is generally preferred. It addresses the non-linearity between house size and price more effectively than the linear model, yields a higher generalized  $R^2$ , and provides a more intuitive interpretation of the relationship through constant elasticity. Furthermore, it helps mitigate the heteroscedasticity observed in the original price-level residuals.